

Did the onboarding redesign work, and should we ship it?


Growth / Product / Risk & Compliance readout · user-level A/B test, broken and re-run

The redesigned onboarding **increases Day-7 Activation by +4.58 pp** — but the first attempt to measure it **broke in production on day 11**, and the clean re-run's guardrail read is **not an unconditional pass**.

Recommendation: ship a monitored rollout with a fraud circuit-breaker — not an unconditional launch. The activation lift is real and statistically resolved. The flagged-fraud guardrail's non-inferiority is not established at this sample size and needs to keep being watched post-launch.

What happened

Growth/Product redesigned onboarding (account created before KYC, not after) and launched it as a 50/50, user-level A/B test, planned for 6 weeks. On **day 11** (2026-07-16), a routine mobile release reset the on-device variant-assignment cache for a share of already-assigned users and pushed new signups into a broken, region-biased assignment fallback. The standing SRM monitor caught it 2 days later. Truncating the data or excluding only the flagged users were both too underpowered and too residually biased to trust (block 04) — the team restarted clean in a shorter, 4-week window with the statistics re-planned from scratch.



Did the redesign work? — Yes, clearly

Day-7 Activation	Control	Treatment	Lift	95% CI	p
Clean re-run (n=9,254)	45.24%	49.83%	+4.58 pp	[+2.55, +6.62]	1.0e-5


The recomputed MDE for the re-run's smaller sample was 2.92 pp (vs. 2.41 pp the original 6-week design could have resolved) — the deadline cost real precision, but the observed effect clears the bar with room to spare.

Is it safe? — One guardrail passes clean, one does not

Guardrail	Control	Treatment	Δ	95% upper bound	Margin	Verdict
KYC rejection rate	10.48%	10.17%	-0.31 pp	+0.92 pp	+1.5 pp	Pass
Flagged-fraud rate	0.95%	1.31%	+0.36 pp	+0.89 pp	+0.5 pp	Not established

The fraud point estimate is small, but the confidence interval reaches past the agreed safety margin — this is "cannot rule out a small increase," not "fraud got worse." In fintech, that distinction is not a rounding error.

Where it works best



Largest for cold, top-of-funnel traffic (paid_social +6.2 pp, the only individually significant channel at this sample size; influencer +7.3 pp, not significant), smallest for warm referral traffic (+1.5 pp) — consistent with the hypothesis that the old flow's document-upload wall cost the most intent from users who had the least of it to spare. The formal interaction test is not significant ($p = 0.52$) — directionally right, not yet proven.

Next steps. (1) Ship to paid_social first — the channel with both the largest and the most confidently estimated lift. (2) Keep the flagged-fraud rate on a live guardrail dashboard through the rollout, not just at this one read. (3) Feed the SRM monitoring runbook's lesson back into the assignment service's cache-reset handling review, not just this one experiment. (4) Re-run the heterogeneity cut once more post-launch data accrues — the channel story is directionally clear but not yet statistically proven.

Synthetic, seed-generated data (a portfolio project). Unit of randomisation = the user, 50/50 via a feature-flag hash. Every headline number is computed on two tracks (pandas + DuckDB) with a parity assertion — `python data_analysis/parity/run_all.py`. The true treatment effect and the exact production-bug magnitudes live only in `etl/config.py::EFFECT / BUG` and were recovered from the data, not read off. Per-block detail: `data_analysis/findings/01 – 05`.