

The A/B test broke. Here's how I salvaged the read.

A production bug corrupted randomization on day 11 — caught by daily SRM monitoring, not by luck.

Fictional Brazilian neobank • synthetic data • real methodology

Day 11

an unrelated mobile app deploy resets the variant-assignment cache.

Some already-assigned users switch arms mid-journey. New signups fall into a region-biased fallback. The question stopped being "did the redesign work?" and became "can this experiment still be trusted?"

Onboarding redesign — fewer steps, ID upload deferred — as a 50/50 user-level A/B test on Day-7 Activation.

PRIMARY METRIC

Day-7 Activation rate — signup to first funded transaction within 7 days

GUARDRAILS

KYC-rejection rate · flagged-fraud rate (30d) — both one-sided non-inferiority tests

Continuous monitoring — not a one-time pre-flight check.

~135K

SIGNUPS

6

WEEK WINDOW

DAILY

SRM MONITORING

Sample-ratio-mismatch, re-derived independently from the raw assignment log every day. That's what caught the break on day 13.

191

contaminated users — plus a fallback skewing treatment share 76.5% (Southeast) to 35.3% (North)

Truncating or excluding the affected users both left an MDE too wide to trust. Restarting clean was the only zero-bias option.

On the clean re-run, the redesign works — +4.58pp Day-7 Activation ($p \approx 1 \times 10^{-5}$).

But the flagged-fraud guardrail doesn't clear non-inferiority. The KYC guardrail does. One clean pass, one not established.

A monitored rollout with a fraud circuit-breaker — not an unconditional launch.

Full methodology, numbers, and code — link in my Featured section.